

# CS 579X Natural Language processing

## Lecture 3: Preprocessing

Prof. Dr. Forrest Sheng Bao  
Dept. of Computer Science  
Iowa State University  
Ames, IA, USA

September 14, 2022

Add a figure of Rosetta stone

# Text encoding in Computers

- ▶ The written system of a language is called a **script**. Many languages do not have written systems.
- ▶ Script elements, e.g., characters, are **encoded**, usually into integers, for storage and representation in computers.
- ▶ For example, in ASCII 95 means “A”.
- ▶ ASCII – Ever wonder why you cannot use certain characters in usernames or passwords?
- ▶ Not every language uses the Latin alphabet. Encoding matters.
- ▶ Local encodings in Europe and Asia
- ▶ Then the birth of Unicode, e.g., UTF8

```
In [20]: s
Out[20]: 'Universitäten'

In [21]: s.encode()
Out[21]: b'Universit\xc3\xa4ten'

In [22]: s.encode().decode('utf-8')
Out[22]: 'Universitäten'

In [23]: s.encode().decode('latin-1')
Out[23]: 'UniversitÃ¼ten'

In [24]: s.encode().decode('GBK')
Out[24]: 'Universit\u0399ten'
```

Figure 1: UTF8\_Byte.png

# Text encoding in Python3

- ▶ In Python 3, a string is by default in UTF-8.
- ▶ This is a huge change from Python 2 to 3.
- ▶ Python 3 has a separate type for strings, bytes.
- ▶ When opening a file, ensure the encoding.

```
In [1]: with open("BIG5.example", 'r') as f:
...:     print (list(f))
...:
.....
13747-600-ish= 3077
.....
UnicodeDecodeError                                Traceback (most recent call last)
<ipython-input-1-91b059321606> in <module>()
----> 1 with open("BIG5.example", 'r') as f:
      2     print (list(f))
      3

/usr/lib/python3.6/codecs.py in decode(self, input, final)
    319         # decode input (taking the buffer into account)
    320         data = self.buffer + input
--> 321         (result, consumed) = self._buffer_decode(data, self.errors, final)
    322         # keep undecoded input until the next call
    323         self.buffer = data[consumed:]

UnicodeDecodeError: 'utf-8' codec can't decode byte 0xa4 in position 0: invalid start byte

In [2]: with open("BIG5.example", 'r', encoding="BIG5") as f:
...:     print (list(f))
...:
['反攻大陸\n', '消滅共匪\n', '\n']
```

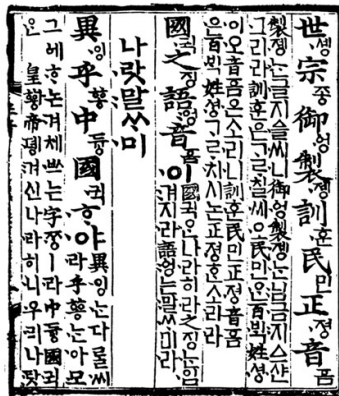
Figure 2: BIG5 encoding

# Unicode can make NLP thorny

- ▶ It simply concatenates different encoding systems. Thus, some symbols are encoded multiple times, there is a cross (U+07d9) in Nko, while there many other crosses in Unicode. How many commas in Unicode?
- ▶ Invisible characters.
- ▶ Potential adversarial attack by using code in different parts of Unicode.
  - ▶ [Bad Characters: Imperceptible NLP Attacks, Boucher et al., 2021](#)
  - ▶ [Hey, AI software developers, you are taking Unicode into account, right . . . right?](#)

# Unicode and CJK I

- ▶ Text is NOT always 1D arrays.
- ▶ Exceptions include Chinese, Korean and Japanese (CJK), where building blocks (strokes, radicals, jamos) are combined in 2D before expanding to 1D arrays.
- ▶ Some papers: Stratos, A Sub-Character Architecture for Korean Language Processing, EMNLP 2017



Ideographic Description Characters<sup>[1][2]</sup>

Official Unicode Consortium code chart<sup>[2]</sup> (PDF)

	0	1	2	3	4	5	6	7	8	9	A	B
U+2FFx	𐄀	𐄁	𐄂	𐄃	𐄄	𐄅	𐄆	𐄇	𐄈	𐄉	𐄊	𐄋

# Unicode and CJK II

- ▶ Hence, encoding characters individually is inefficient. Imaging that you have to encode every word in English using a unique code. You would need  $\log_2 50000$  bits.



- ▶ For CJK, Unicode has 48 strokes (U+31C0..U+31EF), 224 radicals (U+2F00..U+2FDF) and 12 ideographic description characters (U+2FF0..U+2FFF). But it still encodes each character!

## Beyond plain text

Some text data are more than **plain text**, including information about formatting and field/structure.

- ▶ PDF: font and coordinate for each character – no strings
- ▶ Markup languages: HTML, XML (including Microsoft Office XML). Easier but not that easy: e.g., the same type of information is in different tags or tags of different attributes, e.g.,

```
<a href="a.html" id="first_link">Go back</a>  
<a href="b.html" id="second_link">Next</a>
```

- ▶ (The evil) Microsoft formats (prior to Office XML)
- ▶ Dictionary-like data: YAML, JSON
- ▶ Text from charts and text from tables



# Preprocessing steps in NLP

1. Get your data (e.g., crawl, scrape)
2. Extract the part you want (e.g., finding abstracts of papers from )
3. Clean up (e.g., fixing typos)
4. Tokenization
5. Normalization (e.g., stemming/lemmatization, uncasing)

# Corpora (pl. of **corpus**)

- ▶ A corpus is a collection of text data.
- ▶ E.g., all articles on Wikipedia form a corpus.
- ▶ E.g., all news from a newspaper form another corpus.
- ▶ E.g., All ACL papers form a corpus – [Hot off the press](#)
- ▶ Many NLP/DL libraries have made corpora easily accessible:
  - ▶ [NLTK corpus](#) – Salute to the pioneers
  - ▶ [Tensorflow Datasets](#)
  - ▶ [Huggingface Datasets](#)

## Try it out yourself

```
import tensorflow_datasets as tfds

for piece in tfds.load("cnn_dailymail",split="test"):
    print (piece['article'])
    break

for piece in tfds.load("billsum",split="test"):
    print (piece['text'])
    break

for piece in tfds.load("big_patent",split="test"):
    print (piece['description'])
    break
```

## Corpora are diverse

- ▶ In terms of content, domain, representation, etc.
- ▶ Some more examples
  - ▶ <http://jmcauley.ucsd.edu/data/amazon/>
  - ▶ <https://webscope.sandbox.yahoo.com/catalog.php?datatype=l>
  - ▶ <https://www.cs.uic.edu/~liub/FBS/sentiment-analysis.html#datasets>
  - ▶ <https://catalog.ldc.upenn.edu/ldc2008t19>

## A corpus can contain noise – should you clean up?

- ▶ They could contain HTML tags or even unmatched HTML tags
- ▶ They could contain ads in news website data
- ▶ Example: [Cornell Newsroom's human evaluation part](#)

## Constructing a corpus: web and web APIs

- ▶ Two ways to make HTTP requests: GET (data are in the URL) and POST (data is not in the URL)
- ▶ Command line tools `wget` or `curl`. Language bindings: e.g., Python's `urllib` (official) or `request`.

```
>>> import urllib.request
>>> response = urllib.request.urlopen('https://www.wikidata.org/w/api.php?action=w
>>> response.read()
b'{"entities":{"Q2283":{"type":"item","id":"Q2283","labels":{"en":{"language":"en"
```

- ▶ Set a timer to randomly do it.
- ▶ Do NOT be too frequent.

# Parsing HTML/XML

- ▶ Use beautifulsoup
- ▶ An HTML or XML file is a nested structure.

# Regular expressions (regex)

- ▶ A regular expression defines how to generate a string by drawing characters from alphabets. You can draw characters from different sets at different “steps”, like the example `abc)*(123)?` below.
- ▶ `(a)*` means  $\{ \epsilon, a, aa, aaa, aaaa, \dots \}$ , where  $\epsilon$  is the empty string.
- ▶ `(ab)*` means  $\{ \epsilon, ab, abab, ababab, abababab, \dots \}$
- ▶ `(a|b)*` means  $\{ \epsilon, a, b, ab, ba, aaa, aab, abb, bbb, bba, baa, bab, \dots \}$  where `|` means “or”.
- ▶ `(a)+` means  $\{ a, aa, aaa, \dots \}$  because `+` means repeating at least once.
- ▶ `(ab)?` means  $\{ ab, \epsilon \}$  because `?` means once or none.
- ▶ `(abc) * (123)?` means  $\{ \epsilon, abc, abc123, abcabc, abcabc123, \dots \}$  any string that begins with the string `abc` and ends with or without the string `123`.
- ▶ `0|(1|2|3|4|5|6|7|8|9)(0|1|2|3|4|5|6|7|8|9)+` means any natural number in common writing format. Note that it does not allow an integer to start with 0.



## Sentence segmentation

- ▶ Breaking a string into sentences.
- ▶ Also called sentence tokenization, although tokenization usually mean word tokenization.
- ▶ Easy: Just split strings based on punctuations
- ▶ However, you may run into redundant punctuations in Unicode.
- ▶ Hard: dot or period? – Solution: write regex rules.
- ▶ Can be done using rules or neurally.

## (word) Tokenization

- ▶ Breaking a string into tokens, not necessarily words.
- ▶ `"This is a class about NLP".split()`
- ▶ Lots of corner cases:
  - ▶ "he isn't happy"
  - ▶ "We saw 2 people in the park after 3pm"
  - ▶ "Welcome to Columbia University in the City of New York."
- ▶ A classic one

## Modern tokenizers

- ▶ WordPiece tokenizer – made popular by BERT
- ▶ Nearly all HuggingFace models have a tokenizer component – because they may encode tokens to different integers
- ▶ Huggingface tokenizers have two version. `fast`, written in Rust, is available sometimes.

# Tokenization examples in NLTK

```
In [1]: import nltk
```

```
In [2]: nltk.tokenize.word_tokenize \  
...: ("I am happy. mr. Wang is happy")
```

```
Out[2]: ['I', 'am', 'happy', '.', 'mr.', 'Wang', 'is', \  
         'happy']
```

```
In [3]: nltk.tokenize.sent_tokenize \  
...: ("I am happy. mr. Wang is happy")
```

```
Out[3]: ['I am happy.', 'mr. Wang is also happy']
```

```
In [4]: nltk.tokenize.sent_tokenize \  
...: ("I am happy, mr. wang is happy")
```

```
Out[4]: ['I am happy, mr. wang is also happy']
```

Many other varieties at <https://www.nltk.org/api/nltk.tokenize.html>

## Tokenization in SpaCy

SpaCy enforces a pipeline approach. SpaCy's tokenization is rule-based.

```
In [32]: nlp=spacy.load("en_core_web_sm", \
...:   exclude=["tok2vec", 'tagger', 'parser', 'ner', 'attribute_ruler', 'lemmatizer']
...:   nlp.add_pipe("sentencizer")
```

```
Out[32]: <spacy.pipeline.sentencizer.Sentencizer at 0x7f487e80fdc0>
```

```
In [33]: [[sent.text for sent in doc.sents ]\
...:       for doc in nlp.pipe(\
...:   ['today is monday. it is the first day.', 'soup is yammy. pizza sucks.'])]
```

```
Out[33]: [['today is monday.', 'it is the first day.'],
          ['soup is yammy.', 'pizza sucks.']]
```

```
In [34]: [[word.text for word in doc ]
...:       for doc in nlp.pipe(\
...:   ['today is monday. it is the first day.', 'soup is yammy. pizza sucks.'])]
```

```
Out[34]: [['today', 'is', 'monday', '.', 'it', 'is', 'the', 'first', 'day', '.'],
          ['soup', 'is', 'yammy', '.', 'pizza', 'sucks', '.']]
```

# Tokenization in Stanza

Stanza, made by Stanford NLP group, is model-based.

```
In [1]: import stanza
```

```
In [2]: nlp = stanza.Pipeline(lang='en', processors='tokenize')
```

```
In [7]: sentences = 'i am very very happy. the weather is very good.'
```

```
In [8]: [[token.text for token in sentence.tokens] for sentence in nlp(sentences).sentences]
```

```
Out[8]:
```

```
[['i', 'am', 'very', 'very', 'happy', '.'],  
 ['the', 'weather', 'is', 'very', 'good', '.']]
```

```
In [9]: [sentence.text for sentence in nlp(sentences).sentences]
```

```
Out[9]: ['i am very very happy.', 'the weather is very good.']
```

## Word segmentation in other languages

- ▶ More complex than in English
- ▶ Chinese
- ▶ Japanese (very complicated, Kanji + Kana)
- ▶ Korean (Hanja, the chinese character part)
- ▶ Thai
- ▶ Vietnamese
- ▶ Hold on, what about emoji?

# Stemming and Lemmatization

- ▶ The English language has inflexions, *inflectional morphology*. We modify words for different grammatical purposes. “I give him a book” vs. “He gave me a book.”
- ▶ The goal of stemming or lemmatization is to map morphologically related tokens to one canonical form, e.g., “him” to “he”.
- ▶ **analytical languages**: little or no inflectional morphology
  - ▶ **isolating languages**: no inflectional morphology
- ▶ **synthetic languages**: lots of morphemes
- ▶ Most mainland, southeast Asian and Oceanic languages (Chinese – especially classical Chinese, Vietnamese, Samoan, etc.) are isolating language.
- ▶ English is an analytical language. Hence, stemming and lemmatization for English is relatively easy.



# Stemming vs lemmatization

- ▶ Two tasks related, but slightly different.
- ▶ Stemming: get the stem, e.g., removing the affix.
- ▶ Lemmatization: get the lemma, the item word in a dictionary.

seek

/sēk/ 

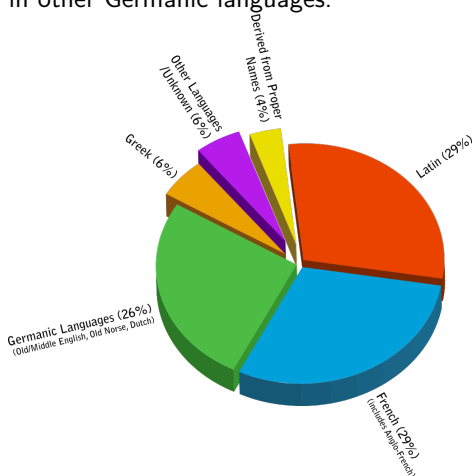
*verb*

verb: **seek**; 3rd person present: **seeks**; past tense: **sought**; past participle: **sought**; gerund or present participle: **seeking**

# The English language is very messed-up

- 1 Ðan ſhe ȝette Old Anȝliſh Tunȝan eode ſpecan
- 2 Than ſhe ȝan to-spaken biſ Middle Englyſſhe Tongue
- 3 Then ſhe wente to ſpake thiſ Early Modern Englyſh Tongue
- 4 Then ſhe went to ſpeak thiſ Late Modern Engliſh Tongue

“English irregular verbs” are actually “regular verbs” in other Germanic languages.



## Building a lemmatizer

- ▶ Take advantage of the dictionary
- ▶ What about words beyond? E.g., `transformative` vs. `transforming`
- ▶ Don't take things for granted: NLTK's "WordNet lemmatizer only removes affixes if the resulting word is in its dictionary" which means that "lying" will not be restored back to "lie".

## Downside of stemming and lemmatization

- ▶ By converting, e.g., “are”, “is”, “am” and “be” into “be”, we lose information that may help use on other tasks.
- ▶ Some combinations are meaningless when tokenized or lemmatized, e.g., in “I should have quited grad school when Sequoia gave me \$2M”, “should have quited” together mean something.
- ▶ Accurate lemmatization requires understanding the role of each word, e.g., the “saw” in “I saw the table (into half)” should be kept instead of being changed to “see.”
- ▶ **Contemporary (2015 forward) NLP does NOT do lemmatization/stemming** because the corpora are large enough.

# Normalization

- ▶ Stemming or lemmatization is under the banner of normalization.
- ▶ Normalization means restoring a word to its canonical form
- ▶ For example, from upper case “University” to lower case “university”.
- ▶ Many NLP models distinguish the cased and uncased one, e.g., BERT's
- ▶ Another example of normalization is dealing with punctuations or escape sequences.
- ▶ And accents

## Stop/background words

- ▶ Some words occur so often in a language that **traditionally** (before Transformer in 2017) their presence does not help the downstream task, such as “is” or “the”.
- ▶ Example: you are building a spam mail filter. You want to count the frequencies of words and use the distribution to predict whether a mail is spam. Both spams and non-spams (hams) contain a lot of “the”.
- ▶ Such words are called stop words or background words.
- ▶ The traditional approach is to remove them, replacing them with spaces.
- ▶ However, the case for keeping is emerging in **contemporary** NLP, thanks to the Transformers. [Ref1](#)  
[Ref2](#)

## Cloud APIs for syntactic analysis

- ▶ Google: <https://cloud.google.com/natural-language/>
- ▶ Microsoft: <https://azure.microsoft.com/en-us/services/cognitive-services/text-analytics/> (no longer has syntactical features)