

# Natural Language Processing

## Lecture IV. Part-of-speech (POS) tagging and Named Entity Recognition (NER)

Forrest Sheng Bao, Ph.D.

Dept. of Computer Science  
Iowa State University  
Ames, IA 50011

September 23, 2021

# Outline

What is POS tagging

POS tagging in HMM

CRF for POS tagging and NER

# POS

Natural	Language	Processing	is	a	field	of	computer	science.
Adj.	n.	n.	v.	dt.	n.	cj.	n.	n.

# Part of speech tagging

- ▶ “In traditional grammar, a part of speech (abbreviated form: PoS or POS) is a category of words (or, more generally, of lexical items) which have similar grammatical properties. ”
- ▶ “In corpus linguistics, part-of-speech tagging (POS tagging or POST), also called grammatical tagging or word-category disambiguation, is the process of marking up a word in a text (corpus) as corresponding to a particular part of speech,
- ▶ based on both its definition and its context, i.e., its relationship with adjacent and related words in a phrase, sentence, or paragraph.”
- ▶ Brill tagger (circa. 1993): the first English POS tagger, rule-based. It assigns initial tags to words first and then use rules to iteratively update tags based on, e.g., context.
- ▶ Brill tagger has hundreds of rules.

# Part of speech tagging

- ▶ “In traditional grammar, a part of speech (abbreviated form: PoS or POS) is a category of words (or, more generally, of lexical items) which have similar grammatical properties. ”
- ▶ “In corpus linguistics, part-of-speech tagging (POS tagging or POST), also called grammatical tagging or word-category disambiguation, is the process of marking up a word in a text (corpus) as corresponding to a particular part of speech,
- ▶ based on both its definition and its context, i.e., its relationship with adjacent and related words in a phrase, sentence, or paragraph.”
- ▶ Brill tagger (circa. 1993): the first English POS tagger, rule-based. It assigns initial tags to words first and then use rules to iteratively update tags based on, e.g., context.
- ▶ Brill tagger has hundreds of rules.

# Part of speech tagging

- ▶ “In traditional grammar, a part of speech (abbreviated form: PoS or POS) is a category of words (or, more generally, of lexical items) which have similar grammatical properties. ”
- ▶ “In corpus linguistics, part-of-speech tagging (POS tagging or POST), also called grammatical tagging or word-category disambiguation, is the process of marking up a word in a text (corpus) as corresponding to a particular part of speech,
- ▶ based on both its definition and its context, i.e., its relationship with adjacent and related words in a phrase, sentence, or paragraph.”
- ▶ Brill tagger (circa. 1993): the first English POS tagger, rule-based. It assigns initial tags to words first and then use rules to iteratively update tags based on, e.g., context.
- ▶ Brill tagger has hundreds of rules.

# Part of speech tagging

- ▶ “In traditional grammar, a part of speech (abbreviated form: PoS or POS) is a category of words (or, more generally, of lexical items) which have similar grammatical properties. ”
- ▶ “In corpus linguistics, part-of-speech tagging (POS tagging or POST), also called grammatical tagging or word-category disambiguation, is the process of marking up a word in a text (corpus) as corresponding to a particular part of speech,
- ▶ based on both its definition and its context, i.e., its relationship with adjacent and related words in a phrase, sentence, or paragraph.”
- ▶ Brill tagger (circa. 1993): the first English POS tagger, rule-based. It assigns initial tags to words first and then use rules to iteratively update tags based on, e.g., context.
- ▶ Brill tagger has hundreds of rules.

# Part of speech tagging

- ▶ “In traditional grammar, a part of speech (abbreviated form: PoS or POS) is a category of words (or, more generally, of lexical items) which have similar grammatical properties. ”
- ▶ “In corpus linguistics, part-of-speech tagging (POS tagging or POST), also called grammatical tagging or word-category disambiguation, is the process of marking up a word in a text (corpus) as corresponding to a particular part of speech,
- ▶ based on both its definition and its context, i.e., its relationship with adjacent and related words in a phrase, sentence, or paragraph.”
- ▶ Brill tagger (circa. 1993): the first English POS tagger, rule-based. It assigns initial tags to words first and then use rules to iteratively update tags based on, e.g., context.
- ▶ Brill tagger has hundreds of rules.



# Tags used in Penn Treebank

- ▶ Nine common parts of speech in English: noun, verb, article, adjective, preposition, pronoun, adverb, conjunction, and interjection.
- ▶ Most NLP researchers use Penn Treebank tags, which are finer than common English POSes mentioned above.
- ▶ [https://www.ling.upenn.edu/courses/Fall\\_2003/ling001/penn\\_treebank\\_pos.html](https://www.ling.upenn.edu/courses/Fall_2003/ling001/penn_treebank_pos.html)

# Tags used in Penn Treebank

- ▶ Nine common parts of speech in English: noun, verb, article, adjective, preposition, pronoun, adverb, conjunction, and interjection.
- ▶ Most NLP researchers use Penn Treebank tags, which are finer than common English POSes mentioned above.
- ▶ [https://www.ling.upenn.edu/courses/Fall\\_2003/ling001/penn\\_treebank\\_pos.html](https://www.ling.upenn.edu/courses/Fall_2003/ling001/penn_treebank_pos.html)

# Tags used in Penn Treebank

- ▶ Nine common parts of speech in English: noun, verb, article, adjective, preposition, pronoun, adverb, conjunction, and interjection.
- ▶ Most NLP researchers use Penn Treebank tags, which are finer than common English POSes mentioned above.
- ▶ [https://www.ling.upenn.edu/courses/Fall\\_2003/ling001/penn\\_treebank\\_pos.html](https://www.ling.upenn.edu/courses/Fall_2003/ling001/penn_treebank_pos.html)

# Probabilistic Model for Tagging

- ▶ Problem of using rule-based system: Very difficult to verify and to scale.
- ▶ Probabilistic approach: there are many sequences of tags, but only one yields (i.e.,  $\text{argmax}$ ) the highest probability.

	Natural	Language	Processing	is	a	field	Of	computer	science.
$T_1$	aj.	n.	n.	v.	dt.	n.	cj.	n.	n.
$T_2$	n.	n.	n.	v.	n.	n.	cj.	n.	n.
$T_3$	v.	n.	av.	v.	dt.	n.	cj.	n.	n.
$T_4$	dt.	n.	n.	v.	v.	n.	aj.	v.	n.
$T_5$	cj.	n.	n.	v.	dt.	n.	cj.	n.	n.

$T_2 - T_5$  apparently make no sense and hence their  $P()$ 's are very low.

- ▶ The goal is to find the most likely sequence of tags ( $T$ ), given the sequence of words ( $W$ ), i.e.,

$$\underset{T}{\text{argmax}} P(T|W)$$

# Probabilistic Model for Tagging

- ▶ Problem of using rule-based system: Very difficult to verify and to scale.
- ▶ Probabilistic approach: there are many sequences of tags, but only one yields (i.e.,  $\text{argmax}$ ) the highest probability.

	Natural	Language	Processing	is	a	field	of	computer	science.
$T_1$	aj.	n.	n.	v.	dt.	n.	cj.	n.	n.
$T_2$	n.	n.	n.	v.	n.	n.	cj.	n.	n.
$T_3$	v.	n.	av.	v.	dt.	n.	cj.	n.	n.
$T_4$	dt.	n.	n.	v.	v.	n.	aj.	v.	n.
$T_5$	cj.	n.	n.	v.	dt.	n.	cj.	n.	n.

$T_2 - T_5$  apparently make no sense and hence their  $P()$ 's are very low.

- ▶ The goal is to find the most likely sequence of tags ( $T$ ), given the sequence of words ( $W$ ), i.e.,

$$\underset{T}{\text{argmax}} P(T|W)$$

# Probabilistic Model for Tagging

- ▶ Problem of using rule-based system: Very difficult to verify and to scale.
- ▶ Probabilistic approach: there are many sequences of tags, but only one yields (i.e.,  $\text{argmax}$ ) the highest probability.

	Natural	Language	Processing	is	a	field	of	computer	science.
$T_1$	aj.	n.	n.	v.	dt.	n.	cj.	n.	n.
$T_2$	n.	n.	n.	v.	n.	n.	cj.	n.	n.
$T_3$	v.	n.	av.	v.	dt.	n.	cj.	n.	n.
$T_4$	dt.	n.	n.	v.	v.	n.	aj.	v.	n.
$T_5$	cj.	n.	n.	v.	dt.	n.	cj.	n.	n.

$T_2 - T_5$  apparently make no sense and hence their  $P()$ 's are very low.

- ▶ The goal is to find the most likely sequence of tags ( $T$ ), given the sequence of words ( $W$ ), i.e.,

$$\underset{T}{\text{argmax}} P(T|W)$$

# A generative model for POS tagging

- ▶ Make use of Bayes' rule:

$$\operatorname{argmax}_{\mathbf{T}} P(\mathbf{T}|\mathbf{W}) = \operatorname{argmax}_{\mathbf{T}} \frac{P(\mathbf{W}|\mathbf{T})P(\mathbf{T})}{P(\mathbf{W})} \quad (1)$$

$$= \operatorname{argmax}_{\mathbf{T}} P(\mathbf{W}|\mathbf{T})P(\mathbf{T}) \quad (2)$$

Note that we go from Eq. (1) to Eq. (2) because  $P(\mathbf{W})$  is not a function of  $\mathbf{T}$ .

- ▶ tag-to-tag *transition* probabilities:  $P(\mathbf{T}) = \prod_{i=1}^{l+1} P_T(t_i|t_{i-1})$ : e.g., NN comes after DET
- ▶ tag-to-word *emission* probabilities:  $P(\mathbf{W}|\mathbf{T}) = \prod_{i=1}^l P_E(w_i|t_i)$ : e.g., “natural” is probably a JJ.
- ▶ A sequence of words is generated in two phases:
  - 1. generate a sequence of tags  $t_1, \dots, t_l$  by sampling from the transition probabilities  $P_T(t_i|t_{i-1})$
  - 2. generate a sequence of words  $w_1, \dots, w_l$  by sampling from the emission probabilities  $P_E(w_i|t_i)$
- ▶ Yes, a semantically meaningless sentence can be tagged and further parsed correctly, e.g., “NLP is a automobile of GOP”.

# A generative model for POS tagging

- ▶ Make use of Bayes' rule:

$$\operatorname{argmax}_{\mathbf{T}} P(\mathbf{T}|\mathbf{W}) = \operatorname{argmax}_{\mathbf{T}} \frac{P(\mathbf{W}|\mathbf{T})P(\mathbf{T})}{P(\mathbf{W})} \quad (1)$$

$$= \operatorname{argmax}_{\mathbf{T}} P(\mathbf{W}|\mathbf{T})P(\mathbf{T}) \quad (2)$$

Note that we go from Eq. (1) to Eq. (2) because  $P(\mathbf{W})$  is not a function of  $\mathbf{T}$ .

- ▶ tag-to-tag *transition* probabilities:  $P(\mathbf{T}) = \prod_{i=1}^{l+1} P_T(t_i|t_{i-1})$ : e.g., NN comes after DET
- ▶ tag-to-word *emission* probabilities:  $P(\mathbf{W}|\mathbf{T}) = \prod_{i=1}^l P_E(w_i|t_i)$ : e.g., “natural” is probably a JJ.
- ▶ A sequence of words is generated in two phases:

- ▶ Yes, a semantically meaningless sentence can be tagged and further parsed correctly, e.g., “NLP is a automobile of GOP”.



# A generative model for POS tagging

- ▶ Make use of Bayes' rule:

$$\operatorname{argmax}_{\mathbf{T}} P(\mathbf{T}|\mathbf{W}) = \operatorname{argmax}_{\mathbf{T}} \frac{P(\mathbf{W}|\mathbf{T})P(\mathbf{T})}{P(\mathbf{W})} \quad (1)$$

$$= \operatorname{argmax}_{\mathbf{T}} P(\mathbf{W}|\mathbf{T})P(\mathbf{T}) \quad (2)$$

Note that we go from Eq. (1) to Eq. (2) because  $P(\mathbf{W})$  is not a function of  $\mathbf{T}$ .

- ▶ tag-to-tag *transition* probabilities:  $P(\mathbf{T}) = \prod_{i=1}^{l+1} P_T(t_i|t_{i-1})$ : e.g., NN comes after DET
- ▶ tag-to-word *emission* probabilities:  $P(\mathbf{W}|\mathbf{T}) = \prod_{i=1}^l P_E(w_i|t_i)$ : e.g., “natural” is probably a JJ.
- ▶ A sequence of words is generated in two phases:
  1. Produce a sequence of tags, e.g., >, NN, DET, ..., based on probability between each two consecutive tags.
- ▶ Yes, a semantically meaningless sentence can be tagged and further parsed correctly, e.g., “NLP is a automobile of GOP”.

# A generative model for POS tagging

- ▶ Make use of Bayes' rule:

$$\operatorname{argmax}_{\mathbf{T}} P(\mathbf{T}|\mathbf{W}) = \operatorname{argmax}_{\mathbf{T}} \frac{P(\mathbf{W}|\mathbf{T})P(\mathbf{T})}{P(\mathbf{W})} \quad (1)$$

$$= \operatorname{argmax}_{\mathbf{T}} P(\mathbf{W}|\mathbf{T})P(\mathbf{T}) \quad (2)$$

Note that we go from Eq. (1) to Eq. (2) because  $P(\mathbf{W})$  is not a function of  $\mathbf{T}$ .

- ▶ tag-to-tag *transition* probabilities:  $P(\mathbf{T}) = \prod_{i=1}^{l+1} P_T(t_i|t_{i-1})$ : e.g., NN comes after DET
- ▶ tag-to-word *emission* probabilities:  $P(\mathbf{W}|\mathbf{T}) = \prod_{i=1}^l P_E(w_i|t_i)$ : e.g., “natural” is probably a JJ.
- ▶ A sequence of words is generated in two phases:
  1. Produce a sequence of tags, e.g.,  $\triangleright$ , NN, DET, ..., based on probability between each two consecutive tags.
  2. For each tag, produce a word, e.g., NN  $\rightarrow$  “language”, DT  $\rightarrow$  “the”.
- ▶ Yes, a semantically meaningless sentence can be tagged and further parsed correctly, e.g., “NLP is a automobile of GOP”.

# A generative model for POS tagging

- ▶ Make use of Bayes' rule:

$$\operatorname{argmax}_{\mathbf{T}} P(\mathbf{T}|\mathbf{W}) = \operatorname{argmax}_{\mathbf{T}} \frac{P(\mathbf{W}|\mathbf{T})P(\mathbf{T})}{P(\mathbf{W})} \quad (1)$$

$$= \operatorname{argmax}_{\mathbf{T}} P(\mathbf{W}|\mathbf{T})P(\mathbf{T}) \quad (2)$$

Note that we go from Eq. (1) to Eq. (2) because  $P(\mathbf{W})$  is not a function of  $\mathbf{T}$ .

- ▶ tag-to-tag *transition* probabilities:  $P(\mathbf{T}) = \prod_{i=1}^{l+1} P_T(t_i|t_{i-1})$ : e.g., NN comes after DET
- ▶ tag-to-word *emission* probabilities:  $P(\mathbf{W}|\mathbf{T}) = \prod_{i=1}^l P_E(w_i|t_i)$ : e.g., “natural” is probably a JJ.
- ▶ A sequence of words is generated in two phases:
  1. Produce a sequence of tags, e.g.,  $\triangleright$ , NN, DET, ..., based on probability between each two consecutive tags.
  2. For each tag, produce a word, e.g., NN  $\rightarrow$  “language”, DT  $\rightarrow$  “the”.
- ▶ Yes, a semantically meaningless sentence can be tagged and further parsed correctly, e.g., “NLP is a automobile of GOP”.

# A generative model for POS tagging

- ▶ Make use of Bayes' rule:

$$\operatorname{argmax}_{\mathbf{T}} P(\mathbf{T}|\mathbf{W}) = \operatorname{argmax}_{\mathbf{T}} \frac{P(\mathbf{W}|\mathbf{T})P(\mathbf{T})}{P(\mathbf{W})} \quad (1)$$

$$= \operatorname{argmax}_{\mathbf{T}} P(\mathbf{W}|\mathbf{T})P(\mathbf{T}) \quad (2)$$

Note that we go from Eq. (1) to Eq. (2) because  $P(\mathbf{W})$  is not a function of  $\mathbf{T}$ .

- ▶ tag-to-tag *transition* probabilities:  $P(\mathbf{T}) = \prod_{i=1}^{l+1} P_T(t_i|t_{i-1})$ : e.g., NN comes after DET
- ▶ tag-to-word *emission* probabilities:  $P(\mathbf{W}|\mathbf{T}) = \prod_{i=1}^l P_E(w_i|t_i)$ : e.g., “natural” is probably a JJ.
- ▶ A sequence of words is generated in two phases:
  1. Produce a sequence of tags, e.g.,  $\triangleright$ , NN, DET, ..., based on probability between each two consecutive tags.
  2. For each tag, produce a word, e.g., NN  $\rightarrow$  “language”, DT  $\rightarrow$  “the”.
- ▶ Yes, a semantically meaningless sentence can be tagged and further parsed correctly, e.g., “NLP is a automobile of GOP”.

# A generative model for POS tagging

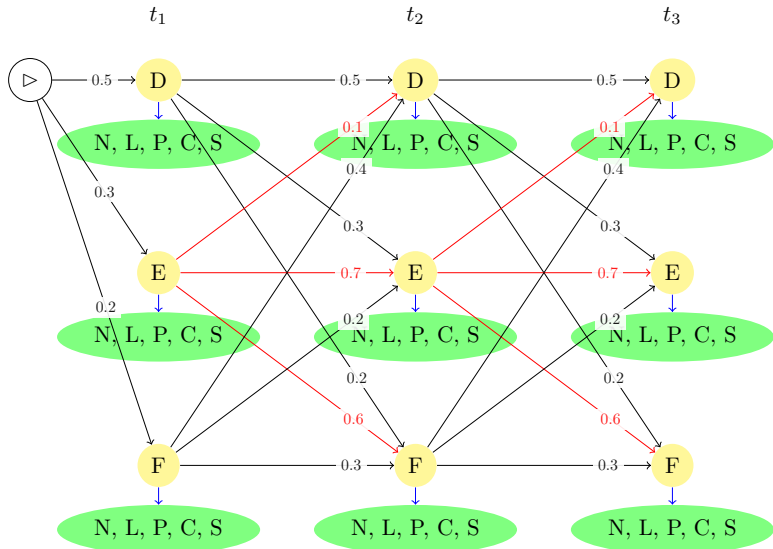
- ▶ Make use of Bayes' rule:

$$\operatorname{argmax}_{\mathbf{T}} P(\mathbf{T}|\mathbf{W}) = \operatorname{argmax}_{\mathbf{T}} \frac{P(\mathbf{W}|\mathbf{T})P(\mathbf{T})}{P(\mathbf{W})} \quad (1)$$

$$= \operatorname{argmax}_{\mathbf{T}} P(\mathbf{W}|\mathbf{T})P(\mathbf{T}) \quad (2)$$

Note that we go from Eq. (1) to Eq. (2) because  $P(\mathbf{W})$  is not a function of  $\mathbf{T}$ .

- ▶ tag-to-tag *transition* probabilities:  $P(\mathbf{T}) = \prod_{i=1}^{l+1} P_T(t_i|t_{i-1})$ : e.g., NN comes after DET
- ▶ tag-to-word *emission* probabilities:  $P(\mathbf{W}|\mathbf{T}) = \prod_{i=1}^l P_E(w_i|t_i)$ : e.g., “natural” is probably a JJ.
- ▶ A sequence of words is generated in two phases:
  1. Produce a sequence of tags, e.g.,  $\triangleright$ , NN, DET, ..., based on probability between each two consecutive tags.
  2. For each tag, produce a word, e.g., NN  $\rightarrow$  “language”, DT  $\rightarrow$  “the”.
- ▶ Yes, a semantically meaningless sentence can be tagged and further parsed correctly, e.g., “NLP is a automobile of GOP”.



Three tags/states: D, E, and F. Five words/observations: N, L, P, C, and S.

# A generative model for POS tagging

<b>W</b>	Natural	Language	Processing	is	a	field	of	computer	science.
<b>T<sub>1</sub></b>	JJ	NN	NN	VBZ	DT	NN	IN	NN	NN.

- ▶ The generative model we just see is a typical Hidden Markov Model (HMM), where a tag is a *state* and a word is an *observation*.
- ▶ In each state/tag, a word/observation emits. After each emission, transit to the next state/tag and emit a word/observation again.
- ▶ Why Markovian? The probability of a tag is only conditioned on the previous tag. No further history. First-order Markovian.
- ▶ For example, the probability of the first tag sequence:  
$$P(\mathbf{T}_1) = P_T(JJ|\triangleright) \times P_T(NN|JJ) \times P_T(NN|NN) \times \dots$$
- ▶ Also, the probability for generating the sentence from the first tag sequence:  $P(\mathbf{W}|\mathbf{T}_1) =$   
$$P_E(\text{"natural"}|JJ) \times P_E(\text{"language"}|NN) \times P_E(\text{"processing"}|NN) \times \dots$$
- ▶ The transition and emission probabilities can be obtained by scanning the corpus once.

# A generative model for POS tagging

- ▶ In principle, we just need to enumerate all possible tag sequences,  $\mathbf{T}_1, \mathbf{T}_2, \dots$  and find the one that yields the largest  $P(\mathbf{W}|\mathbf{T})P(\mathbf{T})$ .
- ▶ But this is costly: If we have  $N$  different tags and  $l$  words in the sentence, there are  $N^l$  possible/tag sequences.
- ▶ Smarter way? Viterbi algorithm.



# A generative model for POS tagging

- ▶ In principle, we just need to enumerate all possible tag sequences,  $\mathbf{T}_1, \mathbf{T}_2, \dots$  and find the one that yields the largest  $P(\mathbf{W}|\mathbf{T})P(\mathbf{T})$ .
- ▶ But this is costly: If we have  $N$  different tags and  $l$  words in the sentence, there are  $N^l$  possible/hidden tag sequences.
- ▶ Smarter way? Viterbi algorithm.

# A generative model for POS tagging

- ▶ In principle, we just need to enumerate all possible tag sequences,  $\mathbf{T}_1, \mathbf{T}_2, \dots$  and find the one that yields the largest  $P(\mathbf{W}|\mathbf{T})P(\mathbf{T})$ .
- ▶ But this is costly: If we have  $N$  different tags and  $l$  words in the sentence, there are  $N^l$  possible/hidden tag sequences.
- ▶ Smarter way? Viterbi algorithm.

# HMM decoding in Viterbi algorithm

- ▶ The problem of estimating the sequence of hidden states given a sequence of observations is known as *decoding* in HMM.
- ▶ Basic principle (Lemma 1):  $\max(x \cdot y) = \max(x) \cdot y$ , if  $x$  is a real variable and  $y$  is a real constant.
- ▶ In each step  $i$  (except the start and end), we have  $N$  possible states/tags  $t_i$ 's, each of which can come from  $N$  possible  $t_{i-1}$ 's.

# HMM decoding in Viterbi algorithm

- ▶ The problem of estimating the sequence of hidden states given a sequence of observations is known as *decoding* in HMM.
- ▶ Basic principle (Lemma 1):  $\max(x \cdot y) = \max(x) \cdot y$ , if  $x$  is a real variable and  $y$  is a real constant.
- ▶ In each step  $i$  (except the start and end), we have  $N$  possible states/tags  $t_i$ 's, each of which can come from  $N$  possible  $t_{i-1}$ 's.

# HMM decoding in Viterbi algorithm

- ▶ The problem of estimating the sequence of hidden states given a sequence of observations is known as *decoding* in HMM.
- ▶ Basic principle (Lemma 1):  $\max(x \cdot y) = \max(x) \cdot y$ , if  $x$  is a real variable and  $y$  is a real constant.
- ▶ In each step  $i$  (except the start and end), we have  $N$  possible states/tags  $t_i$ 's, each of which can come from  $N$  possible  $t_{i-1}$ 's.

# Viterbi algorithm in math induction

- ▶ If the sentence has only one word:  $\mathbf{W} = [w_1]$ . The best tag  $t_1$  should maximize  $P_T(t_1|\triangleright)P_E(w_1|t_1)$  where  $\triangleright = t_0$  is the beginning of the sentence.
- ▶ If it has two:  $\mathbf{W} = [w_1, w_2]$ . The best tags  $t_1$  and  $t_2$  shall maximize  $\prod_{i=1}^2 P_T(t_i|t_{i-1})P_E(w_i|t_i) = P_T(t_1|\triangleright)P_E(w_1|t_1)P_T(t_2|t_1)P_E(w_2|t_2)$ . Just search over the  $N^2$  combinations of  $t_1$  and  $t_2$ , time complexity  $O(N^2)$ .
- ▶ If it has three:  $\mathbf{W} = [w_1, w_2, w_3]$ . Given tags  $t_2$  and  $t_3$ ,

$$\begin{aligned} & \max P_T(t_1|\triangleright)P_E(w_1|t_1)P_T(t_2|t_1)P_E(w_2|t_2) \quad \overbrace{P_T(t_3|t_2)P_E(w_3|t_3)}^{\text{both constants}} \\ & = [\max P_T(t_1|\triangleright)P_E(w_1|t_1)P_T(t_2|t_1)P_E(w_2|t_2)] \quad P_T(t_3|t_2)P_E(w_3|t_3) \end{aligned}$$

- ▶ No need to check  $N^3$  combinations of tags  $t_1$ ,  $t_2$  and  $t_3$ , many of which will not maximize the final number regardless of the value of  $P_T(t_3|t_2)P_E(w_3|t_3)$ .

# Viterbi algorithm in math induction

- ▶ If the sentence has only one word:  $\mathbf{W} = [w_1]$ . The best tag  $t_1$  should maximize  $P_T(t_1|\triangleright)P_E(w_1|t_1)$  where  $\triangleright = t_0$  is the beginning of the sentence.
- ▶ If it has two:  $\mathbf{W} = [w_1, w_2]$ . The best tags  $t_1$  and  $t_2$  shall maximize  $\prod_{i=1}^2 P_T(t_i|t_{i-1})P_E(w_i|t_i) = P_T(t_1|\triangleright)P_E(w_1|t_1)P_T(t_2|t_1)P_E(w_2|t_2)$ . Just search over the  $N^2$  combinations of  $t_1$  and  $t_2$ , time complexity  $O(N^2)$ .
- ▶ If it has three:  $\mathbf{W} = [w_1, w_2, w_3]$ . Given tags  $t_2$  and  $t_3$ ,

$$\begin{aligned} & \max P_T(t_1|\triangleright)P_E(w_1|t_1)P_T(t_2|t_1)P_E(w_2|t_2) \quad \overbrace{P_T(t_3|t_2)P_E(w_3|t_3)}^{\text{both constants}} \\ &= [\max P_T(t_1|\triangleright)P_E(w_1|t_1)P_T(t_2|t_1)P_E(w_2|t_2)] \quad P_T(t_3|t_2)P_E(w_3|t_3) \end{aligned}$$

- ▶ No need to check  $N^3$  combinations of tags  $t_1$ ,  $t_2$  and  $t_3$ , many of which will not maximize the final number regardless of the value of  $P_T(t_3|t_2)P_E(w_3|t_3)$ .

# Viterbi algorithm in math induction

- ▶ If the sentence has only one word:  $\mathbf{W} = [w_1]$ . The best tag  $t_1$  should maximize  $P_T(t_1|\triangleright)P_E(w_1|t_1)$  where  $\triangleright = t_0$  is the beginning of the sentence.
- ▶ If it has two:  $\mathbf{W} = [w_1, w_2]$ . The best tags  $t_1$  and  $t_2$  shall maximize  $\prod_{i=1}^2 P_T(t_i|t_{i-1})P_E(w_i|t_i) = P_T(t_1|\triangleright)P_E(w_1|t_1)P_T(t_2|t_1)P_E(w_2|t_2)$ . Just search over the  $N^2$  combinations of  $t_1$  and  $t_2$ , time complexity  $O(N^2)$ .
- ▶ If it has three:  $\mathbf{W} = [w_1, w_2, w_3]$ . Given tags  $t_2$  and  $t_3$ ,

$$\begin{aligned} & \max P_T(t_1|\triangleright)P_E(w_1|t_1)P_T(t_2|t_1)P_E(w_2|t_2) \quad \overbrace{P_T(t_3|t_2)P_E(w_3|t_3)}^{\text{both constants}} \\ & = [\max P_T(t_1|\triangleright)P_E(w_1|t_1)P_T(t_2|t_1)P_E(w_2|t_2)] \quad P_T(t_3|t_2)P_E(w_3|t_3) \end{aligned}$$

- ▶ No need to check  $N^3$  combinations of tags  $t_1$ ,  $t_2$  and  $t_3$ , many of which will not maximize the final number regardless of the value of  $P_T(t_3|t_2)P_E(w_3|t_3)$ .



# Viterbi algorithm in math induction

- ▶ If the sentence has only one word:  $\mathbf{W} = [w_1]$ . The best tag  $t_1$  should maximize  $P_T(t_1|\triangleright)P_E(w_1|t_1)$  where  $\triangleright = t_0$  is the beginning of the sentence.
- ▶ If it has two:  $\mathbf{W} = [w_1, w_2]$ . The best tags  $t_1$  and  $t_2$  shall maximize  $\prod_{i=1}^2 P_T(t_i|t_{i-1})P_E(w_i|t_i) = P_T(t_1|\triangleright)P_E(w_1|t_1)P_T(t_2|t_1)P_E(w_2|t_2)$ . Just search over the  $N^2$  combinations of  $t_1$  and  $t_2$ , time complexity  $O(N^2)$ .
- ▶ If it has three:  $\mathbf{W} = [w_1, w_2, w_3]$ . Given tags  $t_2$  and  $t_3$ ,

$$\begin{aligned} & \max P_T(t_1|\triangleright)P_E(w_1|t_1)P_T(t_2|t_1)P_E(w_2|t_2) \quad \overbrace{P_T(t_3|t_2)P_E(w_3|t_3)}^{\text{both constants}} \\ & = [\max P_T(t_1|\triangleright)P_E(w_1|t_1)P_T(t_2|t_1)P_E(w_2|t_2)] \quad P_T(t_3|t_2)P_E(w_3|t_3) \end{aligned}$$

- ▶ No need to check  $N^3$  combinations of tags  $t_1$ ,  $t_2$  and  $t_3$ , many of which will not maximize the final number regardless of the value of  $P_T(t_3|t_2)P_E(w_3|t_3)$ .

# HMM decoding in Viterbi algorithm

- ▶ Let's generalize:

$$\begin{aligned} & \max \Pi P_T(t_i | t_{i-1}) P_E(w_i | t_i) & P_T(t_{i+1} | t_i) P_E(w_{i+1} | t_{i+1}) \\ = & [\max \Pi P_T(t_i | t_{i-1}) P_E(w_i | t_i)] & P_T(t_{i+1} | t_i) P_E(w_{i+1} | t_{i+1}) \end{aligned}$$

- ▶ Denote  $v(\overbrace{i}^{\text{step or word index}}, \overbrace{j}^{\text{state or tag index}})$  as the probability that the HMM is in state  $j$  after seeing the first  $i$  observations and passing through **the most probable** preceding sequence of states. We call  $v(i, j)$  the *previous Viterbi path probability*.



$$v(i, j) = \max_{j=1}^N v(i-1, j) P_T(t_i | t_j) P_E(w_i | t_i)$$

- ▶ If we repeat this step by step, we can find the maximum  $P(\mathbf{T} | \mathbf{W})$ .
- ▶ Then a traceback allows us to find the tags that maximizes it.

# HMM decoding in Viterbi algorithm

- ▶ Let's generalize:

$$\begin{aligned} & \max \Pi P_T(t_i | t_{i-1}) P_E(w_i | t_i) & P_T(t_{i+1} | t_i) P_E(w_{i+1} | t_{i+1}) \\ = & [\max \Pi P_T(t_i | t_{i-1}) P_E(w_i | t_i)] & P_T(t_{i+1} | t_i) P_E(w_{i+1} | t_{i+1}) \end{aligned}$$

- ▶ Denote  $v(\overbrace{i}^{\text{step or word index}}, \overbrace{j}^{\text{state or tag index}})$  as the probability that the HMM is in state  $j$  after seeing the first  $i$  observations and passing through **the most probable** preceding sequence of states. We call  $v(i, j)$  the *previous Viterbi path probability*.

$$v(i, j) = \max_{j=1}^N v(i-1, j) P_T(t_i | t_j) P_E(w_i | t_i)$$

- ▶ If we repeat this step by step, we can find the maximum  $P(\mathbf{T} | \mathbf{W})$ .
- ▶ Then a traceback allows us to find the tags that maximizes it.

# HMM decoding in Viterbi algorithm

- ▶ Let's generalize:

$$\begin{aligned} & \max \Pi P_T(t_i | t_{i-1}) P_E(w_i | t_i) & P_T(t_{i+1} | t_i) P_E(w_{i+1} | t_{i+1}) \\ = & [\max \Pi P_T(t_i | t_{i-1}) P_E(w_i | t_i)] & P_T(t_{i+1} | t_i) P_E(w_{i+1} | t_{i+1}) \end{aligned}$$

- ▶ Denote  $v(\overbrace{i}^{\text{step or word index}}, \overbrace{j}^{\text{state or tag index}})$  as the probability that the HMM is in state  $j$  after seeing the first  $i$  observations and passing through **the most probable** preceding sequence of states. We call  $v(i, j)$  the *previous Viterbi path probability*.



$$v(i, j) = \max_{j=1}^N v(i-1, j) P_T(t_i | t_j) P_E(w_i | t_i)$$

- ▶ If we repeat this step by step, we can find the maximum  $P(\mathbf{T} | \mathbf{W})$ .
- ▶ Then a traceback allows us to find the tags that maximizes it.

# HMM decoding in Viterbi algorithm

- ▶ Let's generalize:

$$\begin{aligned} & \max \Pi P_T(t_i | t_{i-1}) P_E(w_i | t_i) & P_T(t_{i+1} | t_i) P_E(w_{i+1} | t_{i+1}) \\ = & [\max \Pi P_T(t_i | t_{i-1}) P_E(w_i | t_i)] & P_T(t_{i+1} | t_i) P_E(w_{i+1} | t_{i+1}) \end{aligned}$$

- ▶ Denote  $v(\overbrace{i}^{\text{step or word index}}, \overbrace{j}^{\text{state or tag index}})$  as the probability that the HMM is in state  $j$  after seeing the first  $i$  observations and passing through **the most probable** preceding sequence of states. We call  $v(i, j)$  the *previous Viterbi path probability*.



$$v(i, j) = \max_{j=1}^N v(i-1, j) P_T(t_i | t_j) P_E(w_i | t_i)$$

- ▶ If we repeat this step by step, we can find the maximum  $P(\mathbf{T} | \mathbf{W})$ .

- ▶ Then a traceback allows us to find the tags that maximizes it.

# HMM decoding in Viterbi algorithm

- ▶ Let's generalize:

$$\begin{aligned} & \max \Pi P_T(t_i | t_{i-1}) P_E(w_i | t_i) & P_T(t_{i+1} | t_i) P_E(w_{i+1} | t_{i+1}) \\ = & [\max \Pi P_T(t_i | t_{i-1}) P_E(w_i | t_i)] & P_T(t_{i+1} | t_i) P_E(w_{i+1} | t_{i+1}) \end{aligned}$$

- ▶ Denote  $v(\overbrace{i}^{\text{step or word index}}, \overbrace{j}^{\text{state or tag index}})$  as the probability that the HMM is in state  $j$  after seeing the first  $i$  observations and passing through **the most probable** preceding sequence of states. We call  $v(i, j)$  the *previous Viterbi path probability*.



$$v(i, j) = \max_{j=1}^N v(i-1, j) P_T(t_i | t_j) P_E(w_i | t_i)$$

- ▶ If we repeat this step by step, we can find the maximum  $P(\mathbf{T} | \mathbf{W})$ .
- ▶ Then a traceback allows us to find the tags that maximizes it.

# HMM decoding in Viterbi algorithm

- ▶ Time complexity:  $O(lN^2)$  instead of  $O(N^3)$  where  $N$  is the number of tags and  $l$  is the length of the sentence under POS-tagging.
- ▶ For details, read chapter 9.4 of <https://web.stanford.edu/~jurafsky/slp3/9.pdf>.

# What do you need?

- ▶ A set of  $N$  states/tags.
- ▶ A set of  $M$  observations/words.
- ▶ Transition probabilities, from one state/tag to another, usually as an  $N \times N$  matrix
- ▶ Emitting probabilities, from one state/tag to an observation/word, usually as another matrix  $N \times M$ .
- ▶ Initial state/tag probabilities, usually denoted as  $\pi_i$ . But this can be easily resolved by introducing a origin state and the transition probabilities from the origin state to all other states.



# What do you need?

- ▶ A set of  $N$  states/tags.
- ▶ A set of  $M$  observations/words.
- ▶ Transition probabilities, from one state/tag to another, usually as an  $N \times N$  matrix
- ▶ Emitting probabilities, from one state/tag to an observation/word, usually as another matrix  $N \times M$ .
- ▶ Initial state/tag probabilities, usually denoted as  $\pi$ . But this can be easily resolved by introducing a origin state and the transition probabilities from the origin state to all other states.

# What do you need?

- ▶ A set of  $N$  states/tags.
- ▶ A set of  $M$  observations/words.
- ▶ Transition probabilities, from one state/tag to another, usually as an  $N \times N$  matrix
- ▶ Emitting probabilities, from one state/tag to an observation/word, usually as another matrix  $N \times M$ .
- ▶ Initial state/tag probabilities, usually denoted as  $\pi$ . But this can be easily resolved by introducing a origin state and the transition probabilities from the origin state to all other states.

# What do you need?

- ▶ A set of  $N$  states/tags.
- ▶ A set of  $M$  observations/words.
- ▶ Transition probabilities, from one state/tag to another, usually as an  $N \times N$  matrix
- ▶ Emitting probabilities, from one state/tag to an observation/word, usually as another matrix  $N \times M$ .
- ▶ Initial state/tag probabilities, usually denoted as  $\pi$ . But this can be easily resolved by introducing a origin state and the transition probabilities from the origin state to all other states.

# What do you need?

- ▶ A set of  $N$  states/tags.
- ▶ A set of  $M$  observations/words.
- ▶ Transition probabilities, from one state/tag to another, usually as an  $N \times N$  matrix
- ▶ Emitting probabilities, from one state/tag to an observation/word, usually as another matrix  $N \times M$ .
- ▶ Initial state/tag probabilities, usually denoted as  $\pi$ . But this can be easily resolved by introducing a origin state and the transition probabilities from the origin state to all other states.

# Computational problems

- ▶ What is the problem of multiplying a lengthy list of probabilities?
- ▶ Like gradient vanishing, the product becomes very very small.
- ▶ Hence, a solution is to logarithmize all probabilities and use summation rather than multiplication.
- ▶ See Neubig's slide 8 on HMM. <http://www.phontron.com/slides/nlp-programming-en-04-hmm.pdf>

# Computational problems

- ▶ What is the problem of multiplying a lengthy list of probabilities?
- ▶ Like gradient vanishing, the product becomes very very small.
- ▶ Hence, a solution is to logarithmize all probabilities and use summation rather than multiplication.
- ▶ See Neubig's slide 8 on HMM. <http://www.phontron.com/slides/nlp-programming-en-04-hmm.pdf>

# Computational problems

- ▶ What is the problem of multiplying a lengthy list of probabilities?
- ▶ Like gradient vanishing, the product becomes very very small.
- ▶ Hence, a solution is to logarithmize all probabilities and use summation rather than multiplication.
- ▶ See Neubig's slide 8 on HMM. <http://www.phontron.com/slides/nlp-programming-en-04-hmm.pdf>

# Computational problems

- ▶ What is the problem of multiplying a lengthy list of probabilities?
- ▶ Like gradient vanishing, the product becomes very very small.
- ▶ Hence, a solution is to logarithmize all probabilities and use summation rather than multiplication.
- ▶ See Neubig's slide 8 on HMM. <http://www.phontron.com/slides/nlp-programming-en-04-hmm.pdf>



# Conditional Random Fields

- ▶ HMM uses Bayes theorem to find the most likely tag sequences  $\operatorname{argmax}_T P(\mathbf{T}|\mathbf{W})$ .
- ▶ CRFs directly estimates it.
- ▶ It works by evaluating the chances of each tag sequence over all possible tag sequences in a softmax fashion:

$$P(\mathbf{T}|\mathbf{W}) = \frac{\exp\left(\sum_{k=1}^K u_k F_k(\mathbf{W}, \mathbf{T})\right)}{\sum_{\mathbf{T}' \in \mathcal{T}} \exp\left(\sum_{k=1}^K u_k F_k(\mathbf{W}, \mathbf{T}')\right)}$$

- ▶ Then just need to find  $\operatorname{argmax}_{T \in \mathcal{T}} P(T|W)$ .
- ▶  $u_k$  is the weight for the  $k$ -th feature  $F_k$  which is a function of both word sequence and tag sequence.
- ▶ Linear-chain CRF:  $F_k(\mathcal{W}, \mathcal{T}) = \sum_{i=1}^l f_k(w_{i-1}, w_i, \mathcal{T}, i)$  The sum of a function of the word sequence and only the current and previous tags.

# Conditional Random Fields

- ▶ HMM uses Bayes theorem to find the most likely tag sequences  $\operatorname{argmax}_T P(\mathbf{T}|\mathbf{W})$ .
- ▶ CRFs directly estimates it.
- ▶ It works by evaluating the chances of each tag sequence over all possible tag sequences in a softmax fashion:

$$P(\mathbf{T}|\mathbf{W}) = \frac{\exp\left(\sum_{k=1}^K u_k F_k(\mathbf{W}, \mathbf{T})\right)}{\sum_{\mathbf{T}' \in \mathcal{T}} \exp\left(\sum_{k=1}^K u_k F_k(\mathbf{W}, \mathbf{T}')\right)}$$

- ▶ Then just need to find  $\operatorname{argmax}_{T \in \mathcal{T}} P(T|W)$ .
- ▶  $u_k$  is the weight for the  $k$ -th feature  $F_k$  which is a function of both word sequence and tag sequence.
- ▶ Linear-chain CRF:  $F_k(\mathcal{W}, \mathcal{T}) = \sum_{i=1}^l f_k(w_{i-1}, w_i, \mathcal{T}, i)$  The sum of a function of the word sequence and only the current and previous tags.

# Conditional Random Fields

- ▶ HMM uses Bayes theorem to find the most likely tag sequences  $\operatorname{argmax}_T P(\mathbf{T}|\mathbf{W})$ .
- ▶ CRFs directly estimates it.
- ▶ It works by evaluating the chances of each tag sequence over all possible tag sequences in a softmax fashion:

$$P(\mathbf{T}|\mathbf{W}) = \frac{\exp\left(\sum_{k=1}^K u_k F_k(\mathbf{W}, \mathbf{T})\right)}{\sum_{\mathbf{T}' \in \mathcal{T}} \exp\left(\sum_{k=1}^K u_k F_k(\mathbf{W}, \mathbf{T}')\right)}$$

- ▶ Then just need to find  $\operatorname{argmax}_{T \in \mathcal{T}} P(T|W)$ .
- ▶  $u_k$  is the weight for the  $k$ -th feature  $F_k$  which is a function of both word sequence and tag sequence.
- ▶ Linear-chain CRF:  $F_k(\mathcal{W}, \mathcal{T}) = \sum_{i=1}^l f_k(w_{i-1}, w_i, \mathcal{T}, i)$  The sum of a function of the word sequence and only the current and previous tags.

# Conditional Random Fields

- ▶ HMM uses Bayes theorem to find the most likely tag sequences  $\operatorname{argmax}_T P(\mathbf{T}|\mathbf{W})$ .
- ▶ CRFs directly estimates it.
- ▶ It works by evaluating the chances of each tag sequence over all possible tag sequences in a softmax fashion:

$$P(\mathbf{T}|\mathbf{W}) = \frac{\exp\left(\sum_{k=1}^K u_k F_k(\mathbf{W}, \mathbf{T})\right)}{\sum_{\mathbf{T}' \in \mathcal{T}} \exp\left(\sum_{k=1}^K u_k F_k(\mathbf{W}, \mathbf{T}')\right)}$$

- ▶ Then just need to find  $\operatorname{argmax}_{T \in \mathcal{T}} P(T|W)$ .
- ▶  $u_k$  is the weight for the  $k$ -th feature  $F_k$  which is a function of both word sequence and tag sequence.
- ▶ Linear-chain CRF:  $F_k(\mathcal{W}, \mathcal{T}) = \sum_{i=1}^l f_k(w_{i-1}, w_i, \mathcal{T}, i)$  The sum of a function of the word sequence and only the current and previous tags.

# Conditional Random Fields

- ▶ HMM uses Bayes theorem to find the most likely tag sequences  $\operatorname{argmax}_T P(\mathbf{T}|\mathbf{W})$ .
- ▶ CRFs directly estimates it.
- ▶ It works by evaluating the chances of each tag sequence over all possible tag sequences in a softmax fashion:

$$P(\mathbf{T}|\mathbf{W}) = \frac{\exp\left(\sum_{k=1}^K u_k F_k(\mathbf{W}, \mathbf{T})\right)}{\sum_{\mathbf{T}' \in \mathcal{T}} \exp\left(\sum_{k=1}^K u_k F_k(\mathbf{W}, \mathbf{T}')\right)}$$

- ▶ Then just need to find  $\operatorname{argmax}_{T \in \mathcal{T}} P(T|W)$ .
- ▶  $u_k$  is the weight for the  $k$ -th feature  $F_k$  which is a function of both word sequence and tag sequence.
- ▶ Linear-chain CRF:  $F_k(\mathcal{W}, \mathcal{T}) = \sum_{i=1}^l f_k(w_{i-1}, w_i, \mathcal{T}, i)$  The sum of a function of the word sequence and only the current and previous tags.

# Conditional Random Fields

- ▶ HMM uses Bayes theorem to find the most likely tag sequences  $\operatorname{argmax}_T P(\mathbf{T}|\mathbf{W})$ .
- ▶ CRFs directly estimates it.
- ▶ It works by evaluating the chances of each tag sequence over all possible tag sequences in a softmax fashion:

$$P(\mathbf{T}|\mathbf{W}) = \frac{\exp\left(\sum_{k=1}^K u_k F_k(\mathbf{W}, \mathbf{T})\right)}{\sum_{\mathbf{T}' \in \mathcal{T}} \exp\left(\sum_{k=1}^K u_k F_k(\mathbf{W}, \mathbf{T}')\right)}$$

- ▶ Then just need to find  $\operatorname{argmax}_{T \in \mathcal{T}} P(T|W)$ .
- ▶  $u_k$  is the weight for the  $k$ -th feature  $F_k$  which is a function of both word sequence and tag sequence.
- ▶ Linear-chain CRF:  $F_k(\mathcal{W}, \mathcal{T}) = \sum_{i=1}^l f_k(w_{i-1}, w_i, \mathcal{T}, i)$  The sum of a function of the word sequence and only the current and previous tags.

# Features of using CRFs in POS tagging

- ▶ **Manually engineered features**
- ▶ See §8.5.1 of Jurafsky's book.
- ▶ Also see <https://towardsdatascience.com/pos-tagging-using-crfs-ea430c5fb78b>
- ▶ It's now even common to use DL to extract features and then hook up to CRF, e.g., BiLSTM-CRF (NAACL 2016)  
<https://arxiv.org/pdf/1508.01991.pdf>
- ▶ Sklearn-CRFSuite <https://sklearn-crfsuite.readthedocs.io/en/latest/tutorial.html>

# Features of using CRFs in POS tagging

- ▶ Manually engineered features
- ▶ See §8.5.1 of Jurafsky's book.
- ▶ Also see <https://towardsdatascience.com/pos-tagging-using-crfs-ea430c5fb78b>
- ▶ It's now even common to use DL to extract features and then hook up to CRF, e.g., BiLSTM-CRF (NAACL 2016)  
<https://arxiv.org/pdf/1508.01991.pdf>
- ▶ Sklearn-CRFSuite <https://sklearn-crfsuite.readthedocs.io/en/latest/tutorial.html>



# Features of using CRFs in POS tagging

- ▶ Manually engineered features
- ▶ See §8.5.1 of Jurafsky's book.
- ▶ Also see <https://towardsdatascience.com/pos-tagging-using-crfs-ea430c5fb78b>
- ▶ It's now even common to use DL to extract features and then hook up to CRF, e.g., BiLSTM-CRF (NAACL 2016)  
<https://arxiv.org/pdf/1508.01991.pdf>
- ▶ Sklearn-CRFSuite <https://sklearn-crfsuite.readthedocs.io/en/latest/tutorial.html>

# Features of using CRFs in POS tagging

- ▶ Manually engineered features
- ▶ See §8.5.1 of Jurafsky's book.
- ▶ Also see <https://towardsdatascience.com/pos-tagging-using-crfs-ea430c5fb78b>
- ▶ It's now even common to use DL to extract features and then hook up to CRF, e.g., BiLSTM-CRF (NAACL 2016)  
<https://arxiv.org/pdf/1508.01991.pdf>
- ▶ Sklearn-CRFSuite <https://sklearn-crfsuite.readthedocs.io/en/latest/tutorial.html>

# Features of using CRFs in POS tagging

- ▶ Manually engineered features
- ▶ See §8.5.1 of Jurafsky's book.
- ▶ Also see <https://towardsdatascience.com/pos-tagging-using-crfs-ea430c5fb78b>
- ▶ It's now even common to use DL to extract features and then hook up to CRF, e.g., BiLSTM-CRF (NAACL 2016)  
<https://arxiv.org/pdf/1508.01991.pdf>
- ▶ Sklearn-CRFSuite <https://sklearn-crfsuite.readthedocs.io/en/latest/tutorial.html>

# Named Entity Recognition (NER)

- ▶ NEs are proper nouns. It is not rare for them to contain more than one word, e.g., New York City.
- ▶ Common categories of NEs: organization, people, location, etc.
- ▶ Just like POS tagging, NER can be modeled as a tagging problem.
- ▶ Instead of deciding the POS tags, we decide a different kind of tags.
- ▶ A common type of tags used in NER is BIO: begin, inside, and outside. See <https://medium.com/analytics-vidhya/bio-tagged-text-to-original-text-99b05da6664>
- ▶ See also <https://github.com/scofield7419/sequence-labeling-BiLSTM-CRF>
- ▶ It seems that CRF is more widely used in NER than in POS tagging.

# Named Entity Recognition (NER)

- ▶ NEs are proper nouns. It is not rare for them to contain more than one word, e.g., New York City.
- ▶ Common categories of NEs: organization, people, location, etc.
- ▶ Just like POS tagging, NER can be modeled as a tagging problem.
- ▶ Instead of deciding the POS tags, we decide a different kind of tags.
- ▶ A common type of tags used in NER is BIO: begin, inside, and outside. See <https://medium.com/analytics-vidhya/bio-tagged-text-to-original-text-99b05da6664>
- ▶ See also <https://github.com/scofield7419/sequence-labeling-BiLSTM-CRF>
- ▶ It seems that CRF is more widely used in NER than in POS tagging.

# Named Entity Recognition (NER)

- ▶ NEs are proper nouns. It is not rare for them to contain more than one word, e.g., New York City.
- ▶ Common categories of NEs: organization, people, location, etc.
- ▶ Just like POS tagging, NER can be modeled as a tagging problem.
- ▶ Instead of deciding the POS tags, we decide a different kind of tags.
- ▶ A common type of tags used in NER is BIO: begin, inside, and outside. See <https://medium.com/analytics-vidhya/bio-tagged-text-to-original-text-99b05da6664>
- ▶ See also <https://github.com/scofield7419/sequence-labeling-BiLSTM-CRF>
- ▶ It seems that CRF is more widely used in NER than in POS tagging.

# Named Entity Recognition (NER)

- ▶ NEs are proper nouns. It is not rare for them to contain more than one word, e.g., New York City.
- ▶ Common categories of NEs: organization, people, location, etc.
- ▶ Just like POS tagging, NER can be modeled as a tagging problem.
- ▶ Instead of deciding the POS tags, we decide a different kind of tags.
- ▶ A common type of tags used in NER is BIO: begin, inside, and outside. See <https://medium.com/analytics-vidhya/bio-tagged-text-to-original-text-99b05da6664>
- ▶ See also <https://github.com/scofield7419/sequence-labeling-BiLSTM-CRF>
- ▶ It seems that CRF is more widely used in NER than in POS tagging.

# Named Entity Recognition (NER)

- ▶ NEs are proper nouns. It is not rare for them to contain more than one word, e.g., New York City.
- ▶ Common categories of NEs: organization, people, location, etc.
- ▶ Just like POS tagging, NER can be modeled as a tagging problem.
- ▶ Instead of deciding the POS tags, we decide a different kind of tags.
- ▶ A common type of tags used in NER is BIO: begin, inside, and outside. See <https://medium.com/analytics-vidhya/bio-tagged-text-to-original-text-99b05da6664>
- ▶ See also <https://github.com/scofield7419/sequence-labeling-BiLSTM-CRF>
- ▶ It seems that CRF is more widely used in NER than in POS tagging.



# Named Entity Recognition (NER)

- ▶ NEs are proper nouns. It is not rare for them to contain more than one word, e.g., New York City.
- ▶ Common categories of NEs: organization, people, location, etc.
- ▶ Just like POS tagging, NER can be modeled as a tagging problem.
- ▶ Instead of deciding the POS tags, we decide a different kind of tags.
- ▶ A common type of tags used in NER is BIO: begin, inside, and outside. See <https://medium.com/analytics-vidhya/bio-tagged-text-to-original-text-99b05da6664>
- ▶ See also <https://github.com/scofield7419/sequence-labeling-BiLSTM-CRF>
- ▶ It seems that CRF is more widely used in NER than in POS tagging.

# Named Entity Recognition (NER)

- ▶ NEs are proper nouns. It is not rare for them to contain more than one word, e.g., New York City.
- ▶ Common categories of NEs: organization, people, location, etc.
- ▶ Just like POS tagging, NER can be modeled as a tagging problem.
- ▶ Instead of deciding the POS tags, we decide a different kind of tags.
- ▶ A common type of tags used in NER is BIO: begin, inside, and outside. See <https://medium.com/analytics-vidhya/bio-tagged-text-to-original-text-99b05da6664>
- ▶ See also <https://github.com/scofield7419/sequence-labeling-BiLSTM-CRF>
- ▶ It seems that CRF is more widely used in NER than in POS tagging.

# Modern ways?

Neural network-based generative models, e.g., seq2seq.